

```
1 import pandas as pd
```

```
1 # Mount Google Drive
2 from google.colab import drive
3 drive.mount('/content/drive')
```

Mounted at /content/drive

```
1 # Read the file into a DataFrame
2 file_path = '/content/drive/MyDrive/trainlstm.csv' # Update the file path with your file location
3 train_df = pd.read_csv(file_path)
4
5 # Display the first few rows of the DataFrame
6 train_df.head()
```

	id	label	tweet
0	1	0	@user when a father is dysfunctional and is s...
1	2	0	@user @user thanks for #lyft credit i can't us...
2	3	0	bihday your majesty
3	4	0	#model i love u take with u all the time in ...
4	5	0	factsguide: society now #motivation

```
1 train_df['label'].value_counts()
```

```
0    29720
1     2242
Name: label, dtype: int64
```

```
1 !wget http://downloads.cs.stanford.edu/nlp/data/glove.6B.zip
```

```
--2024-03-23 01:42:19-- http://downloads.cs.stanford.edu/nlp/data/glove.6B.zip
Resolving downloads.cs.stanford.edu (downloads.cs.stanford.edu)... 171.64.64.22
Connecting to downloads.cs.stanford.edu (downloads.cs.stanford.edu)|171.64.64.22|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 862182613 (822M) [application/zip]
Saving to: 'glove.6B.zip'
```

```
glove.6B.zip      100%[=====] 822.24M  5.23MB/s   in 2m 39s
```

```
2024-03-23 01:44:59 (5.16 MB/s) - 'glove.6B.zip' saved [862182613/862182613]
```

```
1 !unzip glove.6B.zip
```

```
Archive:  glove.6B.zip
  inflating: glove.6B.50d.txt
  inflating: glove.6B.100d.txt
  inflating: glove.6B.200d.txt
  inflating: glove.6B.300d.txt
```

```
1 import numpy as np
2
3 words = dict()
4
5 def add_to_dict(d, filename):
6     with open(filename, 'r') as f:
7         for line in f.readlines():
8             line = line.split(' ')
9
10            try:
11                d[line[0]] = np.array(line[1:], dtype=float)
12            except:
13                continue
14
15 add_to_dict(words, 'glove.6B.50d.txt')
16 words
```

```

0.57067 , -0.1036 , 0.20422 , 0.078308 , -0.42795 ,
-1.7984 , -0.27865 , 0.11954 , -0.12689 , 0.031744 ,
3.8631 , -0.17786 , -0.082434 , -0.62698 , 0.26497 ,
-0.057185 , -0.073521 , 0.46103 , 0.30862 , 0.12498 ,
-0.48609 , -0.0080272, 0.031184 , -0.36576 , -0.42699 ,
0.42164 , -0.11666 , -0.50703 , -0.027273 , -0.53285 ]),
'a': array([ 0.21705 , 0.46515 , -0.46757 , 0.10082 , 1.0135 , 0.74845 ,
-0.53104 , -0.26256 , 0.16812 , 0.13182 , -0.24909 , -0.44185 ,
-0.21739 , 0.51004 , 0.13448 , -0.43141 , -0.03123 , 0.20674 ,
-0.78138 , -0.20148 , -0.097401, 0.16088 , -0.61836 , -0.18504 ,
-0.12461 , -2.2526 , -0.22321 , 0.5043 , 0.32257 , 0.15313 ,
3.9636 , -0.71365 , -0.67012 , 0.28388 , 0.21738 , 0.14433 ,
0.25926 , 0.23434 , 0.4274 , -0.44451 , 0.13813 , 0.36973 ,
-0.64289 , 0.024142, -0.039315, -0.26037 , 0.12017 , -0.043782,
0.41013 , 0.1796 ]),
"'': array([ 0.25769 , 0.45629 , -0.76974 , -0.37679 , 0.59272 , -0.063527,
0.20545 , -0.57385 , -0.29009 , -0.13662 , 0.32728 , 1.4719 ,
-0.73681 , -0.12036 , 0.71354 , -0.46098 , 0.65248 , 0.48887 ,
-0.51558 , 0.039951, -0.34307 , -0.014087, 0.86488 , 0.3546 ,
0.7999 , -1.4995 , -1.8153 , 0.41128 , 0.23921 , -0.43139 ,
3.6623 , -0.79834 , -0.54538 , 0.16943 , -0.82017 , -0.3461 ,
0.69495 , -1.2256 , -0.17992 , -0.057474, 0.030498, -0.39543 ,
-0.38515 , -1.0002 , 0.087599, -0.31009 , -0.34677 , -0.31438 ,
0.75004 , 0.97065 ]),
"'s": array([ 0.23727 , 0.40478 , -0.20547 , 0.58805 , 0.65533 ,
0.32867 , -0.81964 , -0.23236 , 0.27428 , 0.24265 ,
0.054992 , 0.16296 , -1.2555 , -0.086437 , 0.44536 ,
0.096561 , -0.16519 , 0.058378 , -0.38598 , 0.086977 ,
0.0033869, 0.55095 , -0.77697 , -0.62096 , 0.092948 ,
-2.5685 , -0.67739 , 0.10151 , -0.48643 , -0.057805 ,
3.1859 , -0.017554 , -0.16138 , 0.055486 , -0.25885 ,
-0.33938 , -0.19928 , 0.26049 , 0.10478 , -0.55934 ,
-0.12342 , 0.65961 , -0.51802 , -0.82995 , -0.082739 ,
0.28155 , -0.423 , -0.27378 , -0.007901 , -0.030231 ]),
'for': array([ 0.15272 , 0.36181 , -0.22168 , 0.066051, 0.13029 , 0.37075 ,
-0.75874 , -0.44722 , 0.22563 , 0.10208 , 0.054225, 0.13494 ,
-0.43052 , -0.2134 , 0.56139 , -0.21445 , 0.077974, 0.10137 ,
-0.51306 , -0.40295 , 0.40639 , 0.23309 , 0.20696 , -0.12668 ,
-0.50634 , -1.7131 , 0.077183, -0.39138 , -0.10594 , -0.23743 ,
3.9552 , 0.66596 , -0.61841 , -0.3268 , 0.37021 , 0.25764 ,
0.38977 , 0.27121 , 0.043024, -0.34322 , 0.020339, 0.2142 ,
0.044097, 0.14003 , -0.20079 , 0.074794, -0.36076 , 0.43382 ,
-0.084617, 0.1214 ]),
'-': array([-0.16768 , 1.2151 , 0.49515 , 0.26836 , -0.4585 ,
-0.23311 , -0.52822 , -1.3557 , 0.16098 , 0.37691 ,
-0.92702 , -0.43904 , -1.0634 , 1.028 , 0.0053943,
0.04153 , -0.018638 , -0.55451 , 0.026166 , 0.28066 ,
-0.66245 , 0.23435 , 0.2451 , 0.025668 , -1.0869 ,
-2.844 , -0.51272 , 0.27286 , 0.0071502, 0.033984 ,
3.9084 , 0.52766 , -0.66899 , 1.8238 , 0.43436 ,
-0.30084 , -0.26996 , 0.4394 , 0.69956 , 0.14885 ,
0.029453 , 1.4888 , 0.52361 , 0.099354 , 1.2515 ,
0.099381 , -0.079261 , -0.30862 , 0.30893 , 0.11023 ]),
'that': array([ 0.88387 , -0.14199 , 0.13566 , 0.098682 , 0.51218 ,
0.40132 , -0.47155 , -0.20747 , 0.01062 , 0.12606

```

```
1 len(words)
```

```
400000
```

```
1 import nltk
```

```
2
```

```
3 nltk.download('wordnet')
```

```
[nltk_data] Downloading package wordnet to /root/nltk_data...
True
```

```
1 tokenizer = nltk.RegexpTokenizer(r"\w+")
```

```
2
```

```
3 tokenizer.tokenize('@user when a father is dysfunctional and is')
```

```
['user', 'when', 'a', 'father', 'is', 'dysfunctional', 'and', 'is']
```

```

1 from nltk.stem import WordNetLemmatizer
2
3 lemmatizer = WordNetLemmatizer()
4
5 lemmatizer.lemmatize('feet')
6
7 def message_to_token_list(s):
8     tokens = tokenizer.tokenize(s)
9     lowercased_tokens = [t.lower() for t in tokens]
10    lemmatized_tokens = [lemmatizer.lemmatize(t) for t in lowercased_tokens]
11    useful_tokens = [t for t in lemmatized_tokens if t in words]
12
13    return useful_tokens
14
15 message_to_token_list('@user feet a fathers is dysfunctional and is')

```

```

['user', 'foot', 'a', 'father', 'is', 'dysfunctional', 'and', 'is']

```

```

1 def message_to_word_vectors(message, word_dict=words):
2     processed_list_of_tokens = message_to_token_list(message)
3
4     vectors = []
5
6     for token in processed_list_of_tokens:
7         if token not in word_dict:
8             continue
9
10        token_vector = word_dict[token]
11        vectors.append(token_vector)
12
13    return np.array(vectors, dtype=float)

```

```

1 message_to_word_vectors('@user when a father is dysfunctional and is').shape

```

(8, 50)

```

1 train_df = train_df.sample(frac=1, random_state=1)
2 train_df.reset_index(drop=True, inplace=True)
3
4 split_index_1 = int(len(train_df) * 0.7)
5 split_index_2 = int(len(train_df) * 0.85)
6
7 train_df, val_df, test_df = train_df[:split_index_1], train_df[split_index_1:split_index_2], train_df[split_index_2:]
8
9 len(train_df), len(val_df), len(test_df)

```

(22373, 4794, 4795)

```
1 test_df
```

	id	label	tweet
27167	21271	0	thats how we do it. #homebrewpeeps
27168	26923	0	i havent ate no fast food/ home cooked food in...
27169	8332	0	i finally found a way how to delete old tweets...
27170	10079	0	because i'm happy clap along if you feel like ...
27171	24049	0	bye â□□ repost from @user be ! #kindness #hap...
...
31957	17290	0	remember itð□□□ #lost #empire #dreams #succes...
31958	5193	0	justice has been served #bosmatrrial
31959	12173	0	ive just repurposed this former mustard jar in...
31960	236	0	the happiest baby ive ever knownð□□□ #cute #sm...
31961	29734	0	#ased bull up: you will dominate your bull a...

4795 rows × 3 columns

```

1 def df_to_X_y(dff):
2     y = dff['label'].to_numpy().astype(int)
3
4     all_word_vector_sequences = []
5
6     for message in dff['tweet']:
7         message_as_vector_seq = message_to_word_vectors(message)
8
9         if message_as_vector_seq.shape[0] == 0:
10             message_as_vector_seq = np.zeros(shape=(1, 50))
11
12         all_word_vector_sequences.append(message_as_vector_seq)
13
14     return all_word_vector_sequences, y

```

```

1 X_train, y_train = df_to_X_y(train_df)
2
3 print(len(X_train), len(X_train[0]))

```

22373 13

```

1 print(len(X_train), len(X_train[2]))

```

22373 7

```


1 sequence_lengths = []
2
3 for i in range(len(X_train)):
4     sequence_lengths.append(len(X_train[i]))
5
6 import matplotlib.pyplot as plt
7
8 plt.hist(sequence_lengths)

```

```

(array([3.493e+03, 7.017e+03, 6.723e+03, 3.786e+03, 1.182e+03, 1.610e+02,
        7.000e+00, 0.000e+00, 1.000e+00, 3.000e+00]),
 array([ 1. ,  6.4, 11.8, 17.2, 22.6, 28. , 33.4, 38.8, 44.2, 49.6, 55. ]),
 <BarContainer object of 10 artists>)

```



```

1 pd.Series(sequence_lengths).describe()

```

```

count    22373.000000
mean      12.692308
std        5.929912
min        1.000000
25%        8.000000
50%       12.000000
75%       17.000000
max       55.000000
dtype: float64

```

```

1 from copy import deepcopy
2
3 def pad_X(X, desired_sequence_length=57):
4     X_copy = deepcopy(X)
5
6     for i, x in enumerate(X):
7         x_seq_len = x.shape[0]
8         sequence_length_difference = desired_sequence_length - x_seq_len
9
10    pad = np.zeros(shape=(sequence_length_difference, 50))
11
12    X_copy[i] = np.concatenate([x, pad])
13
14    return np.array(X_copy).astype(float)

```

```
1 X_train = pad_X(X_train)
```

```
2
```

```
3 X_train.shape
```

```
(22373, 57, 50)
```

```
1 y_train.shape
```

```
(22373,)
```

```
1 X_val, y_val = df_to_X_y(val_df)
```

```
2 X_val = pad_X(X_val)
```

```
3
```

```
4 X_val.shape, y_val.shape
```

```
((4794, 57, 50), (4794,))
```

```
1 X_test, y_test = df_to_X_y(test_df)
```

```
2 X_test = pad_X(X_test)
```

```
3
```

```
4 X_test.shape, y_test.shape
```

```
((4795, 57, 50), (4795,))
```

```
1 from tensorflow.keras import layers
```

```
2 from tensorflow.keras.models import Sequential
```

```
3
```

```
4 model = Sequential([])
```

```
5
```

```
6 model.add(layers.Input(shape=(57, 50)))
```

```
7 model.add(layers.LSTM(64, return_sequences=True))
```

```
8 model.add(layers.Dropout(0.2))
```

```
9 model.add(layers.LSTM(64, return_sequences=True))
```

```
10 model.add(layers.Dropout(0.2))
```

```
11 model.add(layers.LSTM(64, return_sequences=True))
```

```
12 model.add(layers.Dropout(0.2))
```

```
13 model.add(layers.Flatten())
```

```
14 model.add(layers.Dense(1, activation='sigmoid'))
```

```
1 model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
=====	=====	=====
lstm (LSTM)	(None, 57, 64)	29440
dropout (Dropout)	(None, 57, 64)	0
lstm_1 (LSTM)	(None, 57, 64)	33024
dropout_1 (Dropout)	(None, 57, 64)	0
lstm_2 (LSTM)	(None, 57, 64)	33024
dropout_2 (Dropout)	(None, 57, 64)	0
flatten (Flatten)	(None, 3648)	0
dense (Dense)	(None, 1)	3649

```
=====
Total params: 99137 (387.25 KB)
Trainable params: 99137 (387.25 KB)
Non-trainable params: 0 (0.00 Byte)
```

```
1 from tensorflow.keras.losses import BinaryCrossentropy
2 from tensorflow.keras.optimizers import Adam
3 from tensorflow.keras.metrics import AUC
4 from tensorflow.keras.callbacks import ModelCheckpoint
5
6 cp = ModelCheckpoint('model/', save_best_only=True)
7
8 model.compile(optimizer=Adam(learning_rate=0.0001),
9               loss=BinaryCrossentropy(),
10              metrics=['accuracy', AUC(name='auc')])
```

```
1 frequencies = pd.value_counts(train_df['label'])
2
3 frequencies
0      20820
1       1553
Name: label, dtype: int64
```

```
1 weights = {0: frequencies.sum() / frequencies[0], 1: frequencies.sum() / frequencies[1]}
2 weights
{0: 1.0745917387127761, 1: 14.406310367031551}
```

```
1 model.fit(X_train, y_train, validation_data=(X_val, y_val), epochs=20, callbacks=[cp], class_weight=weights)
```

```
Epoch 1/20
700/700 [=====] - 34s 34ms/step - loss: 0.9951 - accuracy: 0.7405 - auc: 0.8390 - val_loss: 0.3809 - val_accu
Epoch 2/20
700/700 [=====] - 14s 20ms/step - loss: 0.8243 - accuracy: 0.8142 - auc: 0.8935 - val_loss: 0.4562 - val_accu
Epoch 3/20
700/700 [=====] - 9s 13ms/step - loss: 0.7721 - accuracy: 0.8284 - auc: 0.9072 - val_loss: 0.4308 - val_accu
Epoch 4/20
700/700 [=====] - 13s 18ms/step - loss: 0.7556 - accuracy: 0.8327 - auc: 0.9112 - val_loss: 0.4260 - val_accu
Epoch 5/20
700/700 [=====] - 20s 29ms/step - loss: 0.7337 - accuracy: 0.8335 - auc: 0.9166 - val_loss: 0.3699 - val_accu
Epoch 6/20
700/700 [=====] - 12s 17ms/step - loss: 0.7159 - accuracy: 0.8411 - auc: 0.9207 - val_loss: 0.4470 - val_accu
Epoch 7/20
700/700 [=====] - 9s 13ms/step - loss: 0.6971 - accuracy: 0.8414 - auc: 0.9248 - val_loss: 0.4461 - val_accu
Epoch 8/20
700/700 [=====] - 22s 31ms/step - loss: 0.6765 - accuracy: 0.8475 - auc: 0.9294 - val_loss: 0.3285 - val_accu
Epoch 9/20
700/700 [=====] - 22s 31ms/step - loss: 0.6568 - accuracy: 0.8477 - auc: 0.9334 - val_loss: 0.2733 - val_accu
Epoch 10/20
700/700 [=====] - 11s 15ms/step - loss: 0.6427 - accuracy: 0.8522 - auc: 0.9361 - val_loss: 0.3016 - val_accu
Epoch 11/20
700/700 [=====] - 12s 18ms/step - loss: 0.6301 - accuracy: 0.8511 - auc: 0.9386 - val_loss: 0.3838 - val_accu
Epoch 12/20
700/700 [=====] - 9s 13ms/step - loss: 0.6145 - accuracy: 0.8533 - auc: 0.9413 - val_loss: 0.3653 - val_accu
Epoch 13/20
700/700 [=====] - 13s 19ms/step - loss: 0.6034 - accuracy: 0.8588 - auc: 0.9437 - val_loss: 0.3758 - val_accu
Epoch 14/20
700/700 [=====] - 10s 14ms/step - loss: 0.5843 - accuracy: 0.8582 - auc: 0.9468 - val_loss: 0.3333 - val_accu
Epoch 15/20
700/700 [=====] - 12s 18ms/step - loss: 0.5671 - accuracy: 0.8637 - auc: 0.9499 - val_loss: 0.3322 - val_accu
Epoch 16/20
700/700 [=====] - 11s 16ms/step - loss: 0.5737 - accuracy: 0.8602 - auc: 0.9485 - val_loss: 0.3282 - val_accu
Epoch 17/20
700/700 [=====] - 23s 32ms/step - loss: 0.5479 - accuracy: 0.8632 - auc: 0.9527 - val_loss: 0.2679 - val_accu
Epoch 18/20
700/700 [=====] - 10s 14ms/step - loss: 0.5369 - accuracy: 0.8676 - auc: 0.9546 - val_loss: 0.2940 - val_accu
Epoch 19/20
700/700 [=====] - 13s 19ms/step - loss: 0.5305 - accuracy: 0.8683 - auc: 0.9556 - val_loss: 0.3888 - val_accu
Epoch 20/20
700/700 [=====] - 9s 13ms/step - loss: 0.5081 - accuracy: 0.8715 - auc: 0.9590 - val_loss: 0.4817 - val_accu
<keras.src.callbacks.History at 0x7c9d4836a650>
```

```
1 from tensorflow.keras.models import load_model
2
3 best_model = load_model('model/')

1 test_predictions = (best_model.predict(X_test) > 0.5).astype(int)
2
3 from sklearn.metrics import classification_report
4
5 print(classification_report(y_test, test_predictions))
```

```
150/150 [=====] - 2s 5ms/step
              precision    recall  f1-score   support

     0         0.99      0.88      0.93       4454
     1         0.35      0.85      0.50        341

 accuracy                   0.88       4795
 macro avg              0.67      0.87      0.72       4795
 weighted avg           0.94      0.88      0.90       4795
```

1